

```
print('Hello World')
```

Hello World

Step 1 Data Exploration & Cleaning & Loading

```
#importing Dependencies
```

```
import pandas as pd
```

```
print (pd.__version__)
```

2.2.3

```
from sqlalchemy import create_engine
```

```
#postgresql toolkit
```

```
import psycopg2
```

```
#Load Data Set
```

```
df = pd.read_csv('Walmart.csv', encoding_errors='ignore')
```

```
df.shape
```

```
df.head()
```

	invoice_id	Branch	City	category	unit_price
0	1	WALM003	San Antonio	Health and beauty	\$74.69
1	2	WALM048	Harlingen	Electronic accessories	\$15.28
2	3	WALM067	Haltom City	Home and lifestyle	\$46.33
3	4	WALM064	Bedford	Health and beauty	\$58.22
4	5	WALM013	Irving	Sports and travel	\$86.31

	quantity	date	time	payment_method	rating	profit_margin
0	7.0	05/01/19	13:08:00	Ewallet	9.1	0.48
1	5.0	08/03/19	10:29:00	Cash	9.6	0.48
2	7.0	03/03/19	13:23:00	Credit card	7.4	0.33
3	8.0	27/01/19	20:33:00	Ewallet	8.4	0.33
4	7.0	08/02/19	10:37:00	Ewallet	5.3	0.48

```
#if you want to see any statistic
```

```
df.describe()
```

	invoice_id	quantity	rating	profit_margin
count	10051.000000	10020.000000	10051.000000	10051.000000
mean	5025.741220	2.353493	5.825659	0.393791
std	2901.174372	1.602658	1.763991	0.090669
min	1.000000	1.000000	3.000000	0.180000
25%	2513.500000	1.000000	4.000000	0.330000
50%	5026.000000	2.000000	6.000000	0.330000
75%	7538.500000	3.000000	7.000000	0.480000
max	10000.000000	10.000000	10.000000	0.570000

```
#df.dtypes() OR
```

```
#TO FIND DATA TYPES & OTHER INFORMATION OF COLUMNS
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10051 entries, 0 to 10050
```

```
Data columns (total 11 columns):
```

#	Column	Non-Null Count	Dtype
0	invoice_id	10051 non-null	int64
1	Branch	10051 non-null	object
2	City	10051 non-null	object
3	category	10051 non-null	object
4	unit_price	10020 non-null	object
5	quantity	10020 non-null	float64
6	date	10051 non-null	object
7	time	10051 non-null	object
8	payment_method	10051 non-null	object
9	rating	10051 non-null	float64
10	profit_margin	10051 non-null	float64

```
dtypes: float64(3), int64(1), object(7)
```

```
memory usage: 863.9+ KB
```

```
#TO FIND DUPLICATES
```

```
df.duplicated().sum()
```

```
#REMOVE DUPLICATES
```

```
df.drop_duplicates(inplace=True) #inplace = true means all duplicates  
will be automatically removed
```

```
df.duplicated().sum()
```

```
#TO CHECK HOW MANY ROWS AND COLUMNS ARE THERE
```

```
df.shape
```

```
(10000, 11)
```

```
#CHECK ANY MISSING VALUES
```

```
df.isnull().sum()
```

```
invoice_id      0
Branch          0
City            0
category        0
unit_price     31
quantity       31
date           0
time           0
payment_method  0
rating          0
profit_margin   0
dtype: int64
```

```
#DROPPING ALL ROWS WITH MISSING RECORDS
```

```
df.dropna(inplace= True) #DROP ALL THE NULL VALUES
```

```
df.isnull().sum()
```

```
df.shape
```

```
(9969, 11)
```

```
#CONVERT DATA TYPES
```

```
df.dtypes
```

```
df[unit_price].astype(float)
```

```
-----
-----
```

```
NameError                                Traceback (most recent call
last)
```

```
Cell In[29], line 3
```

```
    1 #CONVERT DATA TYPES
```

```
    2 df.dtypes
```

```
----> 3 df[unit_price].astype(float)
```

```
NameError: name 'unit_price' is not defined
```

```
#df['unit_price'] = df['unit_price'].str.replace('$', '').astype(float)
```

```
df['unit_price'] = df['unit_price'].astype(str).str.replace('$', '',
regex=False).astype(float)
```

```
df.head()
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 9969 entries, 0 to 9999
```

```
Data columns (total 11 columns):
```

```

#      Column      Non-Null Count  Dtype
---  -
0      invoice_id    9969 non-null    int64
1      Branch        9969 non-null    object
2      City          9969 non-null    object
3      category       9969 non-null    object
4      unit_price     9969 non-null    float64
5      quantity       9969 non-null    float64
6      date           9969 non-null    object
7      time           9969 non-null    object
8      payment_method  9969 non-null    object
9      rating         9969 non-null    float64
10     profit_margin   9969 non-null    float64
dtypes: float64(4), int64(1), object(6)
memory usage: 934.6+ KB

```

```

#what are the current columns
df.columns

```

```

#CREATE A COLUMN
df['total'] = df['unit_price'] * df['quantity']

```

```
df.head()
```

```

      invoice_id  Branch      City      category
unit_price \
0              1  WALM003  San Antonio  Health and beauty
74.69
1              2  WALM048  Harlingen  Electronic accessories
15.28
2              3  WALM067  Haltom City  Home and lifestyle
46.33
3              4  WALM064   Bedford   Health and beauty
58.22
4              5  WALM013   Irving   Sports and travel
86.31

      quantity      date      time payment_method  rating  profit_margin
total
0          7.0  05/01/19  13:08:00      Ewallet      9.1          0.48
522.83
1          5.0  08/03/19  10:29:00      Cash      9.6          0.48
76.40
2          7.0  03/03/19  13:23:00  Credit card      7.4          0.33
324.31
3          8.0  27/01/19  20:33:00      Ewallet      8.4          0.33
465.76
4          7.0  08/02/19  10:37:00      Ewallet      5.3          0.48
604.17

```

```
#CONNECT TO POSTGRESQL
#psql
```

```
host = localhost
port = 5432
user = postgres
password = 'password'
```

```
df.shape
```

```
(9969, 12)
```

```
df.to_csv('walmart_clean_data.csv',index=False)
```

```
#CONNECTION:
```

```
help(create_engine)
```

Help on function create_engine in module sqlalchemy.engine.create:

```
create_engine(url: 'Union[str, _url.URL]', **kwargs: 'Any') ->
'Engine'
```

Create a new :class:`_engine.Engine` instance.

The standard calling form is to send the :ref:`URL <database_urls>` as the first positional argument, usually a string that indicates database dialect and connection arguments::

```
engine =
create_engine("postgresql+psycopg2://scott:tiger@localhost/test")
```

.. note::

Please review :ref:`database_urls` for general guidelines in composing URL strings. In particular, special characters, such as those often part of passwords, must be URL encoded to be properly parsed.

Additional keyword arguments may then follow it which establish various options on the resulting :class:`_engine.Engine` and its underlying :class:`.Dialect` and :class:`_pool.Pool` constructs::

```
engine = create_engine(
    "mysql+mysqldb://scott:tiger@hostname/dbname",
    pool_recycle=3600,
    echo=True,
)
```

The string form of the URL is
```dialect[+driver]://user:password@host/dbname[?key=value..]```,  
where  
```dialect``` is a database name such as ```mysql```, ```oracle```,  
```postgresql```, etc., and ```driver``` the name of a DBAPI, such as  
```psycopg2```, ```pyodbc```, ```cx_oracle```, etc. Alternatively,  
the URL can be an instance of `:class:`~sqlalchemy.engine.url.URL``.

```**kwargs``` takes a wide variety of options which are routed  
towards their appropriate components. Arguments may be specific  
to  
the `:class:`~_engine.Engine``, the underlying `:class:`~.Dialect``,  
as well as the  
`:class:`~_pool.Pool``. Specific dialects also accept keyword  
arguments that  
are unique to that dialect. Here, we describe the parameters  
that are common to most `:func:`~_sa.create_engine()`` usage.

Once established, the newly resulting `:class:`~_engine.Engine`` will  
request a connection from the underlying `:class:`~_pool.Pool`` once  
`:meth:`~_engine.Engine.connect`` is called, or a method which  
depends on it  
such as `:meth:`~_engine.Engine.execute`` is invoked. The  
`:class:`~_pool.Pool`` in turn  
will establish the first actual DBAPI connection when this request  
is received. The `:func:`~_sa.create_engine`` call itself does  
`**not**`  
establish any actual DBAPI connections directly.

.. seealso::

- `:doc:`~core/engines``
- `:doc:`~dialects/index``
- `:ref:`~connections_toplevel``

`:param connect_args`: a dictionary of options which will be  
passed directly to the DBAPI's ```connect()``` method as  
additional keyword arguments. See the example  
at `:ref:`~custom_dbapi_args``.

`:param creator`: a callable which returns a DBAPI connection.  
This creation function will be passed to the underlying  
connection pool and will be used to create all new database  
connections. Usage of this function causes connection  
parameters specified in the URL argument to be bypassed.

This hook is not as flexible as the newer  
`:meth:`~_events.DialectEvents.do_connect`` hook which allows

complete control over how a connection is made to the database, given the full set of URL arguments and state beforehand.

.. seealso::

:meth:``\_events.DialectEvents.do\_connect`` - event hook that allows full control over DBAPI connection mechanics.

:ref:``custom\_dbapi\_args``

:param echo=False: if True, the Engine will log all statements as well as a ``repr()`` of their parameter lists to the default log handler, which defaults to ``sys.stdout`` for output. If set to the string ``"debug"`` , result rows will be printed to the standard output as well. The ``echo`` attribute of ``Engine`` can be modified at any time to turn logging on and off; direct control of logging is also available using the standard Python ``logging`` module.

.. seealso::

:ref:``dbengine\_logging`` - further detail on how to configure logging.

:param echo\_pool=False: if True, the connection pool will log informational output such as when connections are invalidated as well as when connections are recycled to the default log handler, which defaults to ``sys.stdout`` for output. If set to the string ``"debug"`` , the logging will include pool checkouts and checkins.

Direct control of logging is also available using the standard Python ``logging`` module.

.. seealso::

:ref:``dbengine\_logging`` - further detail on how to configure logging.

:param empty\_in\_strategy: No longer used; SQLAlchemy now uses "empty set" behavior for IN in all cases.

:param enable\_from\_linting: defaults to True. Will emit a warning if a given SELECT statement is found to have un-linked FROM elements which would cause a cartesian product.

.. versionadded:: 1.4

.. seealso::

:ref:`change\_4737`

:param execution\_options: Dictionary execution options which will be applied to all connections. See :meth:`~sqlalchemy.engine.Connection.execution\_options`

:param future: Use the 2.0 style :class:`\_engine.Engine` and :class:`\_engine.Connection` API.

As of SQLAlchemy 2.0, this parameter is present for backwards compatibility only and must remain at its default value of ``True``.

The :paramref:`\_sa.create\_engine.future` parameter will be deprecated in a subsequent 2.x release and eventually removed.

.. versionadded:: 1.4

.. versionchanged:: 2.0 All :class:`\_engine.Engine` objects are "future" style engines and there is no longer a ``future=False`` mode of operation.

.. seealso::

:ref:`migration\_20\_toplevel`

:param hide\_parameters: Boolean, when set to True, SQL statement parameters will not be displayed in INFO logging nor will they be formatted into the string representation of :class:`.StatementError` objects.

.. versionadded:: 1.3.8

.. seealso::



`:ref:`dbengine_logging`` - further detail on how to configure logging.

`:param implicit_returning=True`: Legacy parameter that may only be set to True. In SQLAlchemy 2.0, this parameter does nothing. In order to disable "implicit returning" for statements invoked by the ORM, configure this on a per-table basis using the `:paramref:`.Table.implicit_returning`` parameter.

`:param insertmanyvalues_page_size`: number of rows to format into an INSERT statement when the statement uses "insertmanyvalues" mode, which is a paged form of bulk insert that is used for many backends when using `:term:`executemany`` execution typically in conjunction with RETURNING. Defaults to 1000, but may also be subject to dialect-specific limiting factors which may override this value on a per-statement basis.

.. versionadded:: 2.0

.. seealso::

`:ref:`engine_insertmanyvalues``

`:ref:`engine_insertmanyvalues_page_size``

`:paramref:`_engine.Connection.execution_options.insertmanyvalues_page_size``

`:param isolation_level`: optional string name of an isolation level which will be set on all new connections unconditionally. Isolation levels are typically some subset of the string names ```"SERIALIZABLE"```, ```"REPEATABLE READ"```, ```"READ COMMITTED"```, ```"READ UNCOMMITTED"``` and ```"AUTOCOMMIT"``` based on backend.

The `:paramref:`_sa.create_engine.isolation_level`` parameter is in contrast to the `:paramref:`.Connection.execution_options.isolation_level`` execution option, which may be set on an individual

`:class:'.Connection'`, as well as the same parameter passed to  
`:meth:'.Engine.execution_options'`, where it may be used to  
create multiple engines with different isolation levels that share a  
common connection pool and dialect.

`.. versionchanged:: 2.0` The  
`:paramref:'.sa.create_engine.isolation_level'`  
parameter has been generalized to work on all dialects  
which support the concept of isolation level, and is provided as a more  
succinct,  
up front configuration switch in contrast to the execution  
option which is more of an ad-hoc programmatic option.

`.. seealso::`

`:ref:'.dbapi_autocommit'`

`:param json_deserializer:` for dialects that support the  
`:class:'.types.JSON'`  
datatype, this is a Python callable that will convert a JSON  
string to a Python object. By default, the Python `''json.loads''`  
function is used.

`.. versionchanged:: 1.3.7` The SQLite dialect renamed this  
from `''_json_deserializer''`.

`:param json_serializer:` for dialects that support  
the `:class:'.types.JSON'`  
datatype, this is a Python callable that will render a given  
object as JSON. By default, the Python `''json.dumps''` function is  
used.

`.. versionchanged:: 1.3.7` The SQLite dialect renamed this  
from `''_json_serializer''`.

`:param label_length=None:` optional integer value which limits  
the size of dynamically generated column labels to that many  
characters. If less than 6, labels are generated as  
`"_(counter)"`. If `''None''`, the value of  
`''dialect.max_identifier_length''`, which may be affected via

the

:paramref:`\_sa.create\_engine.max\_identifier\_length` parameter, is used instead. The value of

:paramref:`\_sa.create\_engine.label\_length` may not be larger than that of

:paramref:`\_sa.create\_engine.max\_identifier\_length`.

.. seealso::

:paramref:`\_sa.create\_engine.max\_identifier\_length`

:param logging\_name: String identifier which will be used within the "name" field of logging records generated within the "sqlalchemy.engine" logger. Defaults to a hexstring of the object's id.

.. seealso::

:ref:`dbengine\_logging` - further detail on how to configure logging.

:paramref:`\_engine.Connection.execution\_options.logging\_to`

ken`

:param max\_identifier\_length: integer; override the max\_identifier\_length determined by the dialect. if ``None`` or zero, has no effect. This is the database's configured maximum number of characters that may be used in a SQL identifier such as a table name, column name, or label name. All dialects determine this value automatically, however in the case of a new database version for which this value has changed but SQLAlchemy's dialect has not been adjusted, the value may be passed here.

.. versionadded:: 1.3.9

.. seealso::

:paramref:`\_sa.create\_engine.label\_length`

:param max\_overflow=10: the number of connections to allow in connection pool "overflow", that is connections that can be opened above and beyond the pool\_size setting, which defaults

to five. this is only used  
with :class:`~sqlalchemy.pool.QueuePool`.

:param module=None: reference to a Python module object (the module itself, not its string name). Specifies an alternate DBAPI module to be used by the engine's dialect. Each sub-dialect references a specific DBAPI which will be imported before first connect. This parameter causes the import to be bypassed, and the given module to be used instead. Can be used for testing of DBAPIs as well as to inject "mock" DBAPI implementations into the :class:`~\_engine.Engine`.

:param paramstyle=None: The ``paramstyle`` <https://legacy.python.org/dev/peps/pep-0249/#paramstyle> to use when rendering bound parameters. This style defaults to the one recommended by the DBAPI itself, which is retrieved from the ``.paramstyle`` attribute of the DBAPI. However, most DBAPIs accept more than one paramstyle, and in particular it may be desirable to change a "named" paramstyle into a "positional" one, or vice versa. When this attribute is passed, it should be one of the values ```"qmark"```, ```"numeric"```, ```"named"```, ```"format"``` or ```"pyformat"```, and should correspond to a parameter style known to be supported by the DBAPI in use.

:param pool=None: an already-constructed instance of :class:`~sqlalchemy.pool.Pool`, such as a :class:`~sqlalchemy.pool.QueuePool` instance. If non-None, this pool will be used directly as the underlying connection pool for the engine, bypassing whatever connection parameters are present in the URL argument. For information on constructing connection pools manually, see :ref:`pooling\_toplevel`.

:param poolclass=None: a :class:`~sqlalchemy.pool.Pool` subclass, which will be used to create a connection pool instance using the connection parameters given in the URL.

Note this differs from ```pool``` in that you don't actually

instantiate the pool in this case, you just indicate what type of pool to be used.

`:param pool_logging_name:` String identifier which will be used within the "name" field of logging records generated within the "sqlalchemy.pool" logger. Defaults to a hexstring of the object's id.

`.. seealso::`

`:ref:`dbengine_logging`` - further detail on how to configure logging.

`:param pool_pre_ping:` boolean, if True will enable the connection pool "pre-ping" feature that tests connections for liveness upon each checkout.

`.. versionadded:: 1.2`

`.. seealso::`

`:ref:`pool_disconnects_pessimistic``

`:param pool_size=5:` the number of connections to keep open inside the connection pool. This used with `:class:`~sqlalchemy.pool.QueuePool`` as well as `:class:`~sqlalchemy.pool.SingletonThreadPool``. With `:class:`~sqlalchemy.pool.QueuePool``, a `pool_size` setting of 0 indicates no limit; to disable pooling, set `poolclass` to `:class:`~sqlalchemy.pool.NullPool`` instead.

`:param pool_recycle=-1:` this setting causes the pool to recycle connections after the given number of seconds has passed. It defaults to -1, or no timeout. For example, setting to 3600 means connections will be recycled after one hour. Note that MySQL in particular will disconnect automatically if no activity is detected on a connection for eight hours (although this is configurable with the MySQLDB connection itself and the server configuration as well).

`.. seealso::`

`:ref:`pool_setting_recycle``

`:param pool_reset_on_return='rollback':` set the  
`:paramref: `_pool.Pool.reset_on_return`` parameter of the  
underlying  
`:class: `_pool.Pool`` object, which can be set to the values  
`"rollback"`, `"commit"`, or `None`.

.. seealso::

`:ref: `pool_reset_on_return``

`:param pool_timeout=30:` number of seconds to wait before giving  
up on getting a connection from the pool. This is only used  
with `:class: `~sqlalchemy.pool.QueuePool``. This can be a float  
but is  
subject to the limitations of Python time functions which may  
not be  
reliable in the tens of milliseconds.

.. note: don't use 30.0 above, it seems to break with  
the `:param` tag

`:param pool_use_lifo=False:` use LIFO (last-in-first-out) when  
retrieving  
connections from `:class: `.QueuePool`` instead of FIFO  
(first-in-first-out). Using LIFO, a server-side timeout scheme  
can  
reduce the number of connections used during non- peak  
periods of  
use. When planning for server-side timeouts, ensure that a  
recycle or  
pre-ping strategy is in use to gracefully handle stale  
connections.

.. versionadded:: 1.3

.. seealso::

`:ref: `pool_use_lifo``

`:ref: `pool_disconnects``

`:param plugins:` string list of plugin names to load. See  
`:class: `.CreateEnginePlugin`` for background.

.. versionadded:: 1.2.3

`:param query_cache_size:` size of the cache used to cache the SQL  
string  
form of queries. Set to zero to disable caching.

The cache is pruned of its least recently used items when its size reaches  $N * 1.5$ . Defaults to 500, meaning the cache will always store at least 500 SQL statements when filled, and will grow up to 750 items at which point it is pruned back down to 500 by removing the 250 least recently used items.

Caching is accomplished on a per-statement basis by generating a cache key that represents the statement's structure, then generating string SQL for the current dialect only if that key is not present in the cache. All statements support caching, however some features such as an INSERT with a large set of parameters will intentionally bypass the cache. SQL logging will indicate statistics for each statement whether or not it were pull from the cache.

.. note:: some ORM functions related to unit-of-work persistence as well as some attribute loading strategies will make use of individual per-mapper caches outside of the main cache.

.. seealso::

:ref:`sql\_caching`

.. versionadded:: 1.4

:param use\_insertmanyvalues: True by default, use the "insertmanyvalues" execution style for INSERT..RETURNING statements by default.

.. versionadded:: 2.0

.. seealso::

:ref:`engine\_insertmanyvalues`

```
#PANDAS data frame to export from dataframe to postgresql
#psql connection
#postgresql+psycopg2://scott:tiger@localhost/test
engine_psql =
create_engine("postgresql+psycopg2://postgres:password@localhost:5432/
```

```
Walmart_DB")
```

```
try:
```

```
 engine_psql
```

```
 print("Connected to PSQL")
```

```
except:
```

```
 print("Unable To Connect")
```

Connected to PSQL

```
df.to_sql(name='Walmart',con=engine_psql,
if_exists='append',index=False)
```

969

```
#To see what all columns are there
```

```
df.columns
```

```
Index(['invoice_id', 'Branch', 'City', 'category', 'unit_price',
 'quantity',
 'date', 'time', 'payment_method', 'rating', 'profit_margin',
 'total'],
 dtype='object')
```

```
#Branch & City the are in upper to convert to lower case
```

```
df.columns = df.columns.str.lower()
```

```
df.columns
```

```
Index(['invoice_id', 'branch', 'city', 'category', 'unit_price',
 'quantity',
 'date', 'time', 'payment_method', 'rating', 'profit_margin',
 'total'],
 dtype='object')
```